

# **Markov Decision Processes**

---

## **Fundamentals of AI (AE1FAI)**

Slides created by Dr Huan Jin  
3/27/2023

# About Me

Huan Jin

- Email: [huan.jin@nottingham.edu.cn](mailto:huan.jin@nottingham.edu.cn)
- Office: PMB438
- Office hour: Tuesday 3:00-4:00pm
- Research interests
  - Artificial Intelligence, Heuristics, Algorithms, Transportation Logistics, Vehicle Routing Problem, Scheduling
  - Web: <https://www.nottingham.edu.cn/en/science-engineering/staffprofile/huan-jin.aspx>

# OUTLINE

- Markov decision processes
- Policy evaluation
- Value iteration

# Markov Decision Process


Question:

How would you get to Dongqian Lake on Saturday afternoon in the least amount of time?

- Bike
- Drive
- Didi
- Subway
- Fly


There is uncertainty in the nature!!

# Review on search problem



State  $s$ , action  $a$   $\xrightarrow{\text{Deterministic}}$  state  $Succ(s, a)$

# Uncertainty in the real word



## Application:

- Robotics: decide where to move, but hit unseen obstacles, etc.
- Resource allocation: decide what to produce, but don't know the customer demand for different products
- Agriculture: decide what to plant, but don't know the weather and crop yield

# Example 1: Dice game

For each round  $r=1,2,\dots$

- You choose **stay** or **quit**.
- If **quit**, you get \$10 and we end the game.
- If **stay**, you get \$4 and then I roll a 6-sided dice.
  - If the dice results in 1 or 2, we end the game.
  - Otherwise, continue to the next round.




Dice:

Reward:

# Policy

A policy is a choice of what action to choose at each state.

If follow policy “stay”:



Expected utility:  $\frac{1}{3}(4) + \frac{2}{3} * \frac{1}{3}(8) + \frac{2}{3} * \frac{2}{3} * \frac{1}{3} * (12) + \dots = ???$



# Policy

If follow policy “quit”:




Expected utility:

$$1(10)=10$$

# MDP for dice game

For each round  $r=1,2,\dots$

- You choose **stay** or **quit**.
- If **quit**, you get \$10 and we end the game.
- If **stay**, you get \$4 and then I roll a 6-sided dice.
  - If the dice results in 1 or 2, we end the game.
  - Otherwise, continue to the next round.



# Markov decision process

A MDP is defined by a tuple  $(S, A, T, R)$ :

$S$ : a set of states

$A$ : a set of actions

$T$ : a transition function,

- $T(s, a, s')$  where  $s \in S$ ,  $a \in A$ ,  $s' \in S$ , sometimes denoted as  $P(s' | s, a)$

$R$ : a reward function,

- $R(s, a, s')$  is reward for the transition  $(s, a, s')$

Sometimes also have

- $\gamma$ : discount factor, ( $0 \leq \gamma \leq 1$ )
- Terminal states: processes end after reaching these states,  $\text{IsEnd}(s) = \text{True}$

# In this example



## Definition: Markov decision process

States: the set of states

$s_{\text{start}} \in \text{States}$ : starting state

$\text{Actions}(s)$ : possible actions from state  $s$


$T(s, a, s')$ : probability of  $s'$  if take action  $a$  in state  $s$

$\text{Reward}(s, a, s')$ : reward for the transition  $(s, a, s')$

$\text{IsEnd}(s)$ : whether at end of game

$0 \leq \gamma \leq 1$ : discount factor (default: 1)

# Search Problem



## Definition: search problem

States: the set of states

$s_{\text{start}} \in \text{States}$ : starting state

$\text{Actions}(s)$ : possible actions from state  $s$

$\text{Succ}(s, a)$ : where we end up if take action  $a$  in state  $s$

$\text{Cost}(s, a)$ : cost for taking action  $a$  in state  $s$

$\text{IsEnd}(s)$ : whether at end

$$\text{Succ}(s, a) \Rightarrow T(s, a, s')$$

$$\text{Cost}(s, a) \Rightarrow \text{Reward}(s, a, s')$$

# Transitions

The transition probabilities  $T(s, a, s')$  specify the probability of ending up in state  $s'$  if taken action  $a$  in state  $s$ .



## Example: transition probabilities

$s$	$a$	$s'$	$T(s, a, s')$
in	quit	end	1
in	stay	in	$2/3$
in	stay	end	$1/3$

# Probabilities sum to 1



## Example: transition probabilities

$s$	$a$	$s'$	$T(s, a, s')$
in	quit	end	1
in	stay	in	$2/3$
in	stay	end	$1/3$


For each state  $s$  and  $a$  :

$$\sum_{s' \in \text{States}} T(s, a, s') = 1$$

# What is a solution?

Search problem: Path (sequence of actions)

MDP:



**Definition: policy**

A **policy**  $\pi$  is a mapping from each state  $s \in \text{States}$  to an action  $a \in \text{Actions}(s)$ .



# Evaluating a policy

## Utility:

- Following a policy yields a random path.
- The **utility** of a policy is the (discounted) sum of the rewards on the path (also a random quantity).

Path	Utility
[in; stay, 4, end]	
[in; stay, 4, in; stay, 4, in; stay, 4, end]	
[in; stay, 4, in; stay, 4, end]	

## Value (expected utility):

- The **value** of a policy is the **expected** utility.

# Discounting

Path:  $s_0, a_1 r_1 s_1, a_2 r_2 s_2, \dots$  (action, reward, new state)

The utility with discount  $\gamma$  is

$$u_1 = r_1 + \gamma r_2 + \gamma^2 r_3 + \gamma^3 r_4 + \dots$$

Discount  $\gamma=1$  (save for the future):

[stay, stay, stay]:  $4+4+4=12$

Discount  $\gamma=0$  (live in the moment):


[stay, stay, stay]:  $4+0+0=4$

Discount  $\gamma=0.5$  (balanced life):

[stay, stay, stay]:  $4+0.5*4+0.5*0.5*4=7$

# Policy evaluation


Plan: define recurrences relating value and Q-value



$$V_\pi(s) = \begin{cases} 0 & \text{if IsEnd}(s) \\ Q_\pi(s, \pi(s)) & \text{otherwise.} \end{cases}$$

$$Q_\pi(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_\pi(s')]$$

# Dice game



Let  $\pi$  be the "stay" policy:  $\pi(\text{in}) = \text{stay}$ .

$$V_{\pi}(\text{end}) = 0$$

$$V_{\pi}(\text{in}) = \frac{1}{3} (4 + V_{\pi}(\text{end})) + \frac{2}{3} (4 + V_{\pi}(\text{in}))$$

In this case, can solve in closed form:

$$V_{\pi}(\text{in}) = \frac{1}{3} 4 + \frac{2}{3} (4 + V_{\pi}(\text{in}))$$

# Policy evaluation

## Iterative algorithm:

Start with arbitrary policy values and repeatedly apply recurrences to converge to true values.

## Algorithms:

Initialize  $V_{\pi}^{(0)}(s) \leftarrow 0$  for all states  $s$ .

For iteration  $t = 1, \dots, T_{PE}$ :

For each state  $s$ :

$$V_{\pi}^{(0)}(s) \leftarrow \underbrace{\sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_{\pi}^{(t-1)}(s')]}_{Q^{(t-1)}(s, \pi(s))}$$

Repeat until :  $\max_{s \in S} |V_{\pi}^{(t)}(s) - V_{\pi}^{(t-1)}(s)| < \varepsilon$

# Policy evaluation on dice game

Let  $\pi$  be the "stay" policy:  $\pi(\text{in}) = \text{stay}$ .

$$V_{\pi}^{(t)}(\text{end}) = 0$$

$$V_{\pi}^{(t)}(\text{in}) = \frac{1}{3} (4 + V_{\pi}^{(t-1)}(\text{end})) + \frac{2}{3} (4 + V_{\pi}^{(t-1)}(\text{in}))$$

$s$	end	in	$(t = 100 \text{ iterations})$
$V_{\pi}^{(t)}$	0.00	12.00	


Converges to  $V_{\pi}(\text{in}) = 12$ .

# Value iteration

**Optimal value:**  $V_{opt}(s)$

The **optimal value** is the maximum value attained by any policy.

# Optimal values and Q-values



Optimal value if take action  $a$  in state  $s$ :


$$Q_{\text{opt}}(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_{\text{opt}}(s')].$$

Optimal value from state  $s$ :

$$V_{\text{opt}}(s) = \begin{cases} 0 & \text{if IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} Q_{\text{opt}}(s, a) & \text{otherwise.} \end{cases}$$



# Optimal policy



Given  $Q_{\text{opt}}$ , read off the optimal policy:

$$\pi_{\text{opt}}(s) = \arg \max_{a \in \text{Actions}(s)} Q_{\text{opt}}(s, a)$$

# Value iteration

## Algorithms:

Initialize  $V_{opt}^{(0)}(s) \leftarrow 0$  for all states  $s$ .

For iteration  $t = 1, \dots, T_{VI}$ :

For each state  $s$ :

$$V_{\pi}^{(0)}(s) \leftarrow \max_{a \in A} \underbrace{\sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_{opt}^{(t-1)}(s')]}_{Q_{opt}^{(t-1)}(s, a)}$$

# Value iteration: dice game

$s$	end	in
$V_{\text{opt}}^{(t)}$	0.00	12.00 ( $t = 100$ iterations)
$\pi_{\text{opt}}(s)$	-	stay

# Example from textbook

Actions succeed with probability 0.8 and move at right angles! with probability 0.1 (remain in the same position when there is a wall). Actions incur a small cost (0.04)."

→	→	→	+1
↑		↑	-1
↑	←	←	←

.812	.868	.912	+1
.762		.660	-1
.705	.655	.611	.388

# Summary

**MDPs** cope with uncertainty.

Solutions are policies rather than paths.

**Policy evaluation** computes policy value (expected utility)

**Value iteration** computes optimal value (maximum expected utility) and optimal policy.