

Operating Systems and Concurrency

Concurrency 1
COMP2007

Dan Marsden
(Geert De Maere)
{Geert.DeMaere,Dan.Marsden}@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2023

Concurrency

Context

- Threads and processes **execute concurrently** or **in parallel** and can **share resources** (e.g., devices, memory – variables and data structures)
 - Multi-programming/multi-processing **improves system utilisation**
- A thread can be **interrupted at any point in time** (timer, I/O)
 - The process state is **saved** in the **process control block**
- The outcome of programs may become **unpredictable**:
 - Sharing data can lead to **inconsistencies** - we can be interrupted part way through doing something.
 - The **outcome of execution** may **depend on the order** in which code gets to run on the CPU.

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```


Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>

int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

```
#include <stdio.h>
#include <pthread.h>



int counter = 0;
void* calc(void* param) {
    int const iterations = 50000000;
    for(int i = 0; i < iterations; i++)
        counter++;
    return 0;
}

int main() {
    pthread_t tid1 = 0, tid2 = 0;
    pthread_create(&tid1, NULL, calc, 0);
    pthread_create(&tid2, NULL, calc, 0);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The value of counter is: %d\n", counter);
}
```

Example

Incrementing a counter

- `counter++` consists of three separate actions:
 - 1 read the value of counter from memory and **store it in a register**
 - 2 add one to the value in the register
 - 3 store the value of the register **in counter** in memory
- The above actions are **NOT** “atomic”¹




¹Icons from <https://www.flaticon.com/>

Example

Incrementing a counter

- `counter++` consists of three separate actions:
 - 1 read the value of counter from memory and **store it in a register**
 - 2 add one to the value in the register
 - 3 store the value of the register **in counter** in memory
- The above actions are **NOT** “atomic”¹





¹Icons from <https://www.flaticon.com/>

Example

Incrementing a counter

- `counter++` consists of three separate actions:
 - 1 read the value of counter from memory and **store it in a register**
 - 2 add one to the value in the register
 - 3 store the value of the register **in counter** in memory
- The above actions are **NOT** “atomic”¹




¹Icons from <https://www.flaticon.com/>

Example

Incrementing a counter

- `counter++` consists of three separate actions:
 - 1 read the value of counter from memory and **store it in a register**
 - 2 add one to the value in the register
 - 3 store the value of the register **in counter** in memory
- The above actions are **NOT** “atomic”¹



¹Icons from <https://www.flaticon.com/>

Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
Add 1 to register value (= 2)  
Store register in counter (= 2)  
...  
...  
...
```

Thread 2:

```
...  
...  
...  
Read counter -> register (= 2)  
Add 1 to register value (= 3)  
Store register in counter (= 3)
```



Example




Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
Add 1 to register value (= 2)  
Store register in counter (= 2)  
...  
...  
...
```

Thread 2:

```
...  
...  
...  
Read counter -> register (= 2)  
Add 1 to register value (= 3)  
Store register in counter (= 3)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
Add 1 to register value (= 2)  
Store register in counter (= 2)  
...  
...  
...
```

Thread 2:

```
...  
...  
...  
Read counter -> register (= 2)  
Add 1 to register value (= 3)  
Store register in counter (= 3)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
Add 1 to register value (= 2)  
Store register in counter (= 2)  
...  
...  
...
```

Thread 2:

```
...  
...  
...  
Read counter -> register (= 2)  
Add 1 to register value (= 3)  
Store register in counter (= 3)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
Add 1 to register value (= 2)  
Store register in counter (= 2)  
...  
...  
...
```

Thread 2:

```
...  
...  
...  
Read counter -> register (= 2)  
Add 1 to register value (= 3)  
Store register in counter (= 3)
```



Example




Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
Add 1 to register value (= 2)  
Store register in counter (= 2)  
...  
...  
...
```

Thread 2:

```
...  
...  
...  
Read counter -> register (= 2)  
Add 1 to register value (= 3)  
Store register in counter (= 3)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
Add 1 to register value (= 2)  
Store register in counter (= 2)  
...  
...  
...
```

Thread 2:

```
...  
...  
...  
Read counter -> register (= 2)  
Add 1 to register value (= 3)  
Store register in counter (= 3)
```



Example



Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
Add 1 to register value (= 2)  
Store register in counter (= 2)  
...  
...  
...
```

Thread 2:

```
...  
...  
...  
Read counter -> register (= 2)  
Add 1 to register value (= 3)  
Store register in counter (= 3)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example




Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example


Incrementing a counter

Thread 1:

```
...  
Read counter -> register (= 1)  
...  
Add 1 register value (= 2)  
Store value in counter (= 2)  
...  
...
```

Thread 2:

```
...  
...  
Read counter -> register (= 1)  
...  
...  
Add 1 to register value (= 2)  
Store register in counter (= 2)
```



Example 2

Shared procedures

- Consider the following **code shared** between threads/processes
- `chin` and `chout` **shared global variables**

```
void print()  
{  
    chin = getchar();  
    chout = chin;  
    putchar(chout);  
}
```

Example 2

Shared procedures

- Consider **two processes/threads** and the following **interleaved sequence of instructions** (they do **NOT** interact):

Thread 1:

...

chin = getchar(); ...

chout = chin;

putchar(chout);

...

...

...

Thread 2:

...

...

...

chin = getchar();

chout = chin;

putchar(chout);

Example 2

Shared procedures

- Consider **two processes/threads** and the following **interleaved sequence of instructions** (they **DO** interact):

Thread 1:

```
...  
chin = getchar(); ...  
...  
chout = chin;  
putchar(chout);  
...  
...
```

Thread 2:

```
...  
chin = getchar();  
...  
...  
chout = chin;  
putchar(chout);
```

Example 3

Bounded Buffers

- Consider a **bounded buffer** in which N **items** can be stored
- A **counter** is maintained to count the number of items currently in the buffer
 - **Incremented** when an item is **added**
 - **Decrement**ed when an item is **removed**
- Similar **concurrency problems** as with the calculation of sums happen when multiple threads read and write to the bounded buffer.

Example 3

Bounded Buffers – Producer/Consumer

```
// producer
while (true) {
    while (counter == BUFFER_SIZE);
    buffer[in] = new_item;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

```
// consumer
while (true) {
    while (counter == 0);
    consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
}
```


Race Conditions

Definition

- Code has a **race condition** if its behaviour is dependent on the timing of when computation is performed.
- A race condition typically occurs when multiple threads **access shared data** and the result is dependent on **the order in which the instructions are interleaved**.
- We will be interested in **mechanisms** to provide **synchronised** access to data and **avoid race conditions**.

Concurrency within the OS

Data Structures

- **Kernels are preemptive** these days
 - **Multiple threads are running** in the kernel.
 - Kernel **code can be interrupted** at any point.
- The kernel maintains **data structures** such as process tables and open file lists:
 - These data structures may be accessed **concurrently**.
 - These can be subject to **concurrency issues**.
- The OS must make sure that interactions within the OS **do not result in race conditions**.

Critical Sections and Mutual Exclusion

- As race conditions can lead to wild, unpredictable and plain wrong behaviour, we want software abstractions help to prevent them.
- A **critical section** is a section of code that can only be run by one thread at a time. This property is referred to as **mutual exclusion**.
- The question then becomes how to enforce mutual exclusion? Potentially:
 - The O/S and compiler could provide **direct support for critical sections** as a programming primitive.
 - The O/S and compiler provide **locks** which can be held by at most one thread at a time.

Mutual Exclusion

Definition

```
do {  
    ...  
    // ENTRY to critical section  
  
    critical section, e.g.counter++;  
  
    // EXIT critical section  
  
    // remaining code  
    ...  
} while (...);
```

Critical Sections, Mutual Exclusion

Definition

- A **solution to the critical section problem** should satisfy the following **requirements**:
 - **Mutual exclusion**: only one process can be in its critical section at any one point in time
 - **Progress**: any process must be able to enter its critical section at some point in time
 - Processes/threads in the “**remaining code**” **do not influence** access to critical sections
 - **Fairness/bounded waiting**: fairly distributed waiting times/processes cannot be made to wait indefinitely
- These requirements have to be satisfied, **independent of the order** in which computations are executed

Enforcing Mutual Exclusion

Approaches

- A standard approach to enforcing mutual exclusion is via locks known as **mutexes**. These can be implemented in various ways:
 - **Software based**: Peterson's solution
 - **Hardware based**: `test_and_set()`, `swap_and_compare()`
 - **O/S based** - the operating system blocks processes waiting for the lock.
- Unfortunately, **mutexes** and other concurrency primitives such as **semaphore** introduce a new problem - **deadlocks**.

Deadlocks

Example

- Assume that X and Y are **mutexes**.
- Thread A and B need to **acquire both mutexes**, and request them in **opposite orders**.
- The following **sequence of events** could occur in a **multi-programmed system**:

```
THREAD A:  
request mutex X  
acquire mutex X  
...  
...  
request mutex Y  
...
```

```
THREAD B:  
...  
...  
request mutex Y  
acquire mutex Y  
...  
request mutex X  
...
```

Deadlocks

Definition

Tanenbaum

*"A set of threads is deadlocked if **each thread** in the set is waiting for **an event** that only the **other thread** in the set can cause"*

- Each **deadlocked thread** is **waiting for** a resource held by **an other deadlocked thread** (which cannot run and hence cannot release the resources)
- This can happen between **any number of threads** and for **any number of resources**



Figure: Deadlocks

Deadlocks

Minimum Conditions

- **Four conditions** must hold for deadlocks to occur (Coffman et al. (1971)):
 - **Mutual exclusion:** a resource can be assigned to at most one process at a time
 - **Hold and wait condition:** a resource can be held whilst requesting new resources
 - **No preemption:** resources cannot be forcefully taken away from a process
 - **Circular wait:** there is a circular chain of two or more processes, waiting for a resource held by the other processes
- **No deadlocks** can occur if one of the conditions is **not satisfied**

Deadlocks

Minimum Conditions

- **Four conditions** must hold for deadlocks to occur (Coffman et al. (1971)):
 - **Mutual exclusion:** a resource can be assigned to at most one process at a time
 - **Hold and wait condition:** a resource can be held whilst requesting new resources
 - **No preemption:** resources cannot be forcefully taken away from a process
 - **Circular wait:** there is a circular chain of two or more processes, waiting for a resource held by the other processes
- **No deadlocks** can occur if one of the conditions is **not satisfied**
- If your **coursework solution deadlocks**, check for the **order in which resources are requested**

Test your understanding

- The code `x != y` doesn't modify anything. Is it certain to occur atomically?
- Can race conditions or deadlocks occur in practice on a machine with a single hardware thread?
- Can two threads running the same function deadlock against each other?