

Operating Systems and Concurrency

Memory Management 3
COMP2007

Geert De Maere
(Dan Marsden)
{Geert.DeMaere,Dan.Marsden}@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2023

Recall


Last Lecture

- **Dynamic relocation & protection, base & bound registers** (**logical** \Rightarrow **physical** address)
- **Dynamic partitioning & segmentation** (**internal** \Rightarrow **external** fragmentation)
- **Free space** management (bitmaps and linked lists)

Sketch of the Scene

Where Did we Get To?

- **Contiguous memory allocation** methods:
 - Fixed (non-)equal size partitions
 - Dynamic partitioning
 - Segmentation (individual segments are contiguous)
 - `malloc()`, `free()`
- The **amount of memory** requested by processes can be **differ**
- Memory will become **fragmented** over time



Sketch of the Scene

What do we Need?

- A method to **keep track** of free memory
- A strategy to **allocate free** memory when requested by a process
- Approaches to **prevent fragmentation** whenever we can

Overview



Goals for Today

- **Memory allocation:**
 - First fit, best fit, next fit, worst fit, quick fit
 - Buddy algorithm
- **Prevent fragmentation - non-contiguous** memory management
 - **(Segmentation)**
 - **Paging** - page tables, address translation

Memory Allocation

First Fit



- **First fit** scans **from the start** of the list until an **sufficiently large gap** is found
 - **if** the space is **the exact size** then all the space is **allocated**
 - **else** the space is **split**:
 - The first entry is set to the **size requested** and marked **“used”**
 - The second entry is set to **remaining size** and marked **“free”**



Memory Allocation

First Fit



- **First fit** scans **from the start** of the list until an **sufficiently large gap** is found
 - **if** the space is **the exact size** then all the space is **allocated**
 - **else** the space is **split**:
 - The first entry is set to the **size requested** and marked **“used”**
 - The second entry is set to **remaining size** and marked **“free”**



Memory Allocation

Next Fit



- The **next fit algorithm** maintains a record of where it got to:
 - It **restarts its search** from **where it stopped last time**
 - It gives an **even chance to all memory** to get allocated (first fit concentrates on the start of the list)
- Simulations have shown that next fit **performs worse** than first fit



Memory Allocation

Next Fit

- The **next fit algorithm** maintains a record of where it got to:
 - It **restarts its search** from **where it stopped last time**
 - It gives an **even chance to all memory** to get allocated (first fit concentrates on the start of the list)
- Simulations have shown that next fit **performs worse** than first fit



Memory Allocation



First vs. Next Fit

- **First fit** is a fast and looks for the **first available hole**
 - It does **not account** for a **better fitting** hole later
 - It may **break up** a **large holes** early in the list
- **Next fit** does not change this

Memory Allocation

Best Fit



- **Best fit** searches the **entire linked** list for the **smallest hole** big enough to satisfy the request
 - It is **slower** than first fit
 - Result in **small leftover holes** (wastes memory)



Memory Allocation

Best Fit



- **Best fit** searches the **entire linked** list for the **smallest hole** big enough to satisfy the request
 - It is **slower** than first fit
 - Result in **small leftover holes** (wastes memory)



Memory Allocation

Worst Fit



- **Best fit** split an empty partition in **small holes**
- **Worst fit** finds the **largest available partition** and splits it
 - The left over part **will still be large** (potentially more useful)
 - Simulations have shown that **worst fit** is **not good** in practice



Memory Allocation

Worst Fit

- **Tiny holes** are created when **best fit** split an empty partition
- The **worst fit algorithm** finds the **largest available empty partition** and splits it
 - The **left over part will still be large** (and **potentially more useful**)
 - Simulations have also shown that worst fit is **not very good either!**



Memory Allocation

Summary

- **First fit**: allocate **first block** that is **large enough**
- **Next fit**: allocate **next block** that is large enough, i.e. **starting from the current location**
- **Best fit**: choose block that **matches** required size **closest** - $O(N)$ complexity
- **Worst fit**: choose the **largest possible block** - $O(N)$ complexity

Dynamic Partitioning



Allocating Available Memory: Quick Fit and Others

- **Quick fit** maintains **lists of commonly used sizes**
 - For example a separate list for each of 4K, 8K, 12K, 16K, etc., holes
 - **Odd sizes** can either go into the **nearest size** or into a **special separate list**
- It is **much faster** to find the required size hole using **quick fit**
- Similar to **best fit**, it has the problem of creating **many tiny holes**
- Finding neighbours for **coalescing** (combining empty partitions) becomes more difficult/time consuming

Memory Allocation

Coalescing



- **Coalescing** (joining together) takes place when two **adjacent entries** in the linked list **become free**
- Both **neighbours** are examined when a **block is freed**
 - if either (or both) are **also free**
 - then the two (or three) **entries are combined** into one larger block by adding up the sizes
 - The earlier block in the linked list gives the **start point**
 - The **separate links are deleted** and a single link inserted



Memory Allocation

Coalescing



- **Coalescing** (joining together) takes place when two **adjacent entries** in the linked list **become free**
- Both **neighbours** are examined when a **block is freed**
 - if either (or both) are **also free**
 - then the two (or three) **entries are combined** into one larger block by adding up the sizes
 - The earlier block in the linked list gives the **start point**
 - The **separate links are deleted** and a single link inserted



Memory Allocation

Coalescing


- **Coalescing** (joining together) takes place when two **adjacent entries** in the linked list **become free**
- Both **neighbours** are examined when a **block is freed**
 - if either (or both) are **also free**
 - then the two (or three) **entries are combined** into one larger block by adding up the sizes
 - The earlier block in the linked list gives the **start point**
 - The **separate links are deleted** and a single link inserted



Memory Allocation

Coalescing

- **Coalescing** (joining together) takes place when two **adjacent entries** in the linked list **become free**
- Both **neighbours** are examined when a **block is freed**
 - if either (or both) are **also free**
 - then the two (or three) **entries are combined** into one larger block by adding up the sizes
 - The earlier block in the linked list gives the **start point**
 - The **separate links are deleted** and a single link inserted



Memory Allocation

Compacting

- Even with coalescing, **free blocks** may still **distributed** across memory
- **Compacting** can be used but is harder to implement and **time consuming**
- Processes may have to be **moved/swapped out**, **free space coalesced**, and processes **swapped back in** at lowest available location

Contiguous Allocation Schemes

Overview and Shortcomings

- Different contiguous memory allocation schemes have different advantages/disadvantages
 - **Mono-programming** is easy but does result in **low resource utilisation**
 - **Fixed partitioning** facilitates **multi-programming** but results in **internal fragmentation**
 - **Dynamic partitioning & segmentation** facilitate **multi-programming**, reduce **internal fragmentation**, but result in **external fragmentation** (allocation methods, coalescing, and compacting help)
- Can we design a memory management scheme that **resolves the shortcomings of contiguous memory schemes**?

Paging

Principles

- **Paging** uses **fixed partitioning** and **code re-location** to devise a new ***non-contiguous*** management scheme:
 - Memory is split into **much smaller blocks**
 - A process is allocated **one or more blocks** (e.g., a 11KB process would take up three 4KB blocks)
 - Blocks do **not** have to be stored **in contiguous in physical memory**
 - The process **perceives** them to be **contiguous**
- Benefits of non-contiguous schemes include:
 - **Internal fragmentation** is reduced to the **last “block”** only
 - **No external fragmentation** (blocks in physical address space are **stacked** directly on top of each other)

Paging

Principles (Cont'ed)




Figure: Paging in physical memory with multiple processes

Paging

Principles (Cont'ed)




Figure: Paging in main memory with multiple processes (Anderson)

Paging

Principles (Cont'ed)




Figure: Concept of Paging (Stallings)

Paging

Principles: Definitions

- A **page** is a small block of **contiguous memory** in the **logical address space**, i.e. as seen by the process
- A **frame** is a **small contiguous block** in **physical memory**
- Pages and frames (commonly) have the **same size**:
 - The size is usually a power of 2
 - Sizes range between 512 bytes and 1Gb

Paging

Relocation

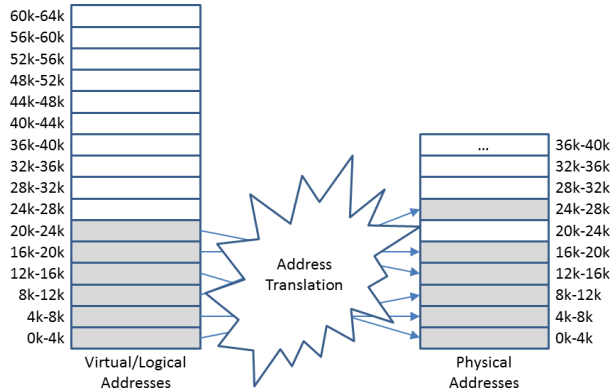


Figure: Address Translation

Segmentation

Segmentation Table




Figure: Logical address for segmentation (offset = position relative to start of segment)

Segment	index	Base	Bound	RWX
Code	00
Data	01
Heap	10
Stack	11

Paging

Page Table



Figure: Logical address for page (offset = position relative to start of segment)

Page Number	Frame Number	RWX
0000	1010
0001	0100
0010	1111
...

Paging

Page Table



Figure: Logical address for page (offset = position relative to start of segment)

Page Number	Frame Number	RWX
0000	1010
0001	0100
0010	1111
...

Paging

Page Table



Figure: Logical address for page (offset = position relative to start of segment)

Page Number	Frame Number	RWX
0000	1010
0001	0100
0010	1111
...

Paging

Page Table



Figure: Logical address for page (offset = position relative to start of segment)

Page Number	Frame Number	RWX
0000	1010
0001	0100
0010	1111
...

Paging

Page Table



Figure: Logical address for page (offset = position relative to start of segment)

Page Number	Frame Number	RWX
0000	1010
0001	0100
0010	1111
...

Paging

Relocation

- **Logical address** (page number, offset within page) needs to be **translated** into a **physical address** (frame number, offset within frame)
- **Multiple “base registers”** will be required:
 - Each **page** has a “**base register**” that identifies the start of the associated **frame**
 - A **set of base registers** must be maintained for every process
- The **set of base registers** is stored in the **page table**

Paging

Relocation: Address Translation

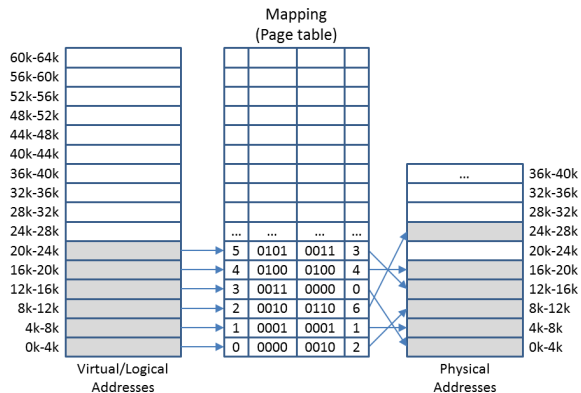


Figure: Address Translation

Paging

Relocation: Page Tables

- The page table maps the **page number** onto the **frame number** (logical \Rightarrow physical address)
 - $\text{frameNumber} = f(\text{pageNumber})$
- The **page number** is used as **index to the page table** that lists the **number of the associated frame**
- From the **frame number** one can calculate the base (offset remains unchanged) is used as **index to the page table** that lists the **number of the associated frame**
- Every process has its **own page table** containing its own **set of “base registers”**
- The **operating system** maintains a **list of free frames**

Paging

Address Translation: Implementation




Figure: Address Translation

Test Your Understanding

Page tables and memory addressing

Page Tables

Given a **64-bit machine** that uses paging, and a page/frame size of **4KB**.

- What would be the maximum **number of frames**? $2^{64} / (4 \times 1024)$
- How many **entries** will we have in the **page table**? $2^{64} / (4 \times 1024)$
- How many **pages** do we have in a **17KB process**?
- How much **memory is wasted** in the last frame?

Recap

Take-Home Message

- Memory **allocation**, **coalescing** and **compacting** in dynamic partitioning
- **Paging**, **page tables**, and **address translation**